

```
!pip install --upgrade pip
```

→ Requirement already satisfied: pip in /usr/local/lib/python3.11/dist-packages (25.2)

```
!pip install imbalanced-learn
```

→ Requirement already satisfied: imbalanced-learn in /usr/local/lib/python3.11/dist-packages (0.13.0)  
 Requirement already satisfied: numpy<3,>=1.24.3 in /usr/local/lib/python3.11/dist-packages (from imbalanced-learn) (2.0.2)  
 Requirement already satisfied: scipy<2,>=1.10.1 in /usr/local/lib/python3.11/dist-packages (from imbalanced-learn) (1.16.1)  
 Requirement already satisfied: scikit-learn<2,>=1.3.2 in /usr/local/lib/python3.11/dist-packages (from imbalanced-learn) (1.6.1)  
 Requirement already satisfied: sklearn-compat<1,>=0.1 in /usr/local/lib/python3.11/dist-packages (from imbalanced-learn) (0.1.3)  
 Requirement already satisfied: joblib<2,>=1.1.1 in /usr/local/lib/python3.11/dist-packages (from imbalanced-learn) (1.5.1)  
 Requirement already satisfied: threadpoolctl<4,>=2.0.0 in /usr/local/lib/python3.11/dist-packages (from imbalanced-learn) (3.6.0)

```
pip install tensorflow
```

→ Requirement already satisfied: tensorflow in /usr/local/lib/python3.11/dist-packages (2.19.0)  
 Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.4.0)  
 Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.6.3)  
 Requirement already satisfied: flatbuffers>=24.3.25 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (25.2.10)  
 Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.6.0)  
 Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.2.0)  
 Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (18.1.1)  
 Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.4.0)  
 Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from tensorflow) (25.0)  
 Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<6.0.0dev,>=3.20.3 in /usr/local/lib/python  
 Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.32.3)  
 Requirement already satisfied: setuptools in /usr/local/lib/python3.11/dist-packages (from tensorflow) (75.2.0)  
 Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.0)  
 Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.1.0)  
 Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (4.14.1)  
 Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.2)  
 Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.74.0)  
 Requirement already satisfied: tensorflow<~2.19.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.19.0)  
 Requirement already satisfied: keras>=3.5.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.10.0)  
 Requirement already satisfied: numpy<2.2.0,>=1.26.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.0.2)  
 Requirement already satisfied: h5py>=3.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.14.0)  
 Requirement already satisfied: ml-dtypes<1.0.0,>=0.5.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.5.3)  
 Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.37.1)  
 Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow)  
 Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.10)  
 Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (2.5  
 Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (202  
 Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.11/dist-packages (from tensorflow<~2.19.0->tensorflow) (3.8.2)  
 Requirement already satisfied: tensorflow-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow<~2.19.  
 Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow<~2.19.0->tensorflow) (3.1.3)  
 Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.11/dist-packages (from astunparse>=1.6.0->tensorflow) (0.45.  
 Requirement already satisfied: rich in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (13.9.4)  
 Requirement already satisfied: namex in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (0.1.0)  
 Requirement already satisfied: optree in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (0.17.0)  
 Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.11/dist-packages (from werkzeug>=1.0.1->tensorflow<~2.19.0->  
 Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tensorflow) (3  
 Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tensorflow)  
 Requirement already satisfied: mdurl~0.1 in /usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0->rich->keras>=3.5.0->te

```
pip install scikeras
```

→ Requirement already satisfied: scikeras in /usr/local/lib/python3.11/dist-packages (0.13.0)  
 Requirement already satisfied: keras>=3.2.0 in /usr/local/lib/python3.11/dist-packages (from scikeras) (3.10.0)  
 Requirement already satisfied: scikit-learn>=1.4.2 in /usr/local/lib/python3.11/dist-packages (from scikeras) (1.6.1)  
 Requirement already satisfied: absl-py in /usr/local/lib/python3.11/dist-packages (from keras>=3.2.0->scikeras) (1.4.0)  
 Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from keras>=3.2.0->scikeras) (2.0.2)  
 Requirement already satisfied: rich in /usr/local/lib/python3.11/dist-packages (from keras>=3.2.0->scikeras) (13.9.4)  
 Requirement already satisfied: namex in /usr/local/lib/python3.11/dist-packages (from keras>=3.2.0->scikeras) (0.1.0)  
 Requirement already satisfied: h5py in /usr/local/lib/python3.11/dist-packages (from keras>=3.2.0->scikeras) (3.14.0)  
 Requirement already satisfied: optree in /usr/local/lib/python3.11/dist-packages (from keras>=3.2.0->scikeras) (0.17.0)  
 Requirement already satisfied: ml-dtypes in /usr/local/lib/python3.11/dist-packages (from keras>=3.2.0->scikeras) (0.5.3)  
 Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from keras>=3.2.0->scikeras) (25.0)  
 Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn>=1.4.2->scikeras) (1.16.1)  
 Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn>=1.4.2->scikeras) (1.5.1)  
 Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn>=1.4.2->scikeras) (3.6  
 Requirement already satisfied: typing-extensions>=4.6.0 in /usr/local/lib/python3.11/dist-packages (from optree->keras>=3.2.0->scikeras)  
 Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.2.0->scikeras) (3.0  
 Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.2.0->scikeras) (2  
 Requirement already satisfied: mdurl~0.1 in /usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0->rich->keras>=3.2.0->sc

```
pip install xgboost
```

```

Requirement already satisfied: xgboost in /usr/local/lib/python3.11/dist-packages (3.0.3)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from xgboost) (2.0.2)
Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.11/dist-packages (from xgboost) (2.23.4)
Requirement already satisfied: scipy in /usr/local/lib/python3.11/dist-packages (from xgboost) (1.16.1)

import csv
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import sklearn
import scipy
import time

from imblearn.over_sampling import SMOTE
from collections import Counter
import itertools

from sklearn.tree import DecisionTreeClassifier
from scipy.stats import pearsonr
from sklearn.ensemble import RandomForestClassifier
from scipy.stats import chi2_contingency
from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.neural_network import MLPClassifier

from sklearn.preprocessing import MinMaxScaler

import tensorflow as tf
from keras.models import Sequential
from keras.layers import LSTM, Dense, Dropout, Activation, BatchNormalization
from scikeras import wrappers
from scikeras.wrappers import KerasClassifier
from sklearn.model_selection import GridSearchCV

from sklearn.metrics import roc_curve, auc as auc_score
from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score

from sklearn.model_selection import cross_val_score, cross_val_predict, cross_validate
from sklearn.metrics import make_scorer, precision_score, recall_score, f1_score, roc_auc_score, accuracy_score
from sklearn.preprocessing import StandardScaler

from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn import svm
from sklearn.svm import SVC
from sklearn.ensemble import AdaBoostClassifier
from sklearn.linear_model import SGDClassifier

import pandas as pd

file_path = "/Cancer_Data.csv"
df = pd.read_csv(file_path)
df.head()

```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10

5 rows × 33 columns

```
df = df.drop(['id'], axis=1)
df.head(5)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	symme
0	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	
1	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	
2	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	
3	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	
4	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	

5 rows × 32 columns

```
df.tail(5)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	syr
564	M	21.56	22.39	142.00	1479.0	0.11100	0.11590	0.24390	0.13890	
565	M	20.13	28.25	131.20	1261.0	0.09780	0.10340	0.14400	0.09791	
566	M	16.60	28.08	108.30	858.1	0.08455	0.10230	0.09251	0.05302	
567	M	20.60	29.33	140.10	1265.0	0.11780	0.27700	0.35140	0.15200	
568	B	7.76	24.54	47.92	181.0	0.05263	0.04362	0.00000	0.00000	

5 rows × 32 columns

## EDA

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

## EDA-END

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

```
from sklearn.preprocessing import LabelEncoder
```

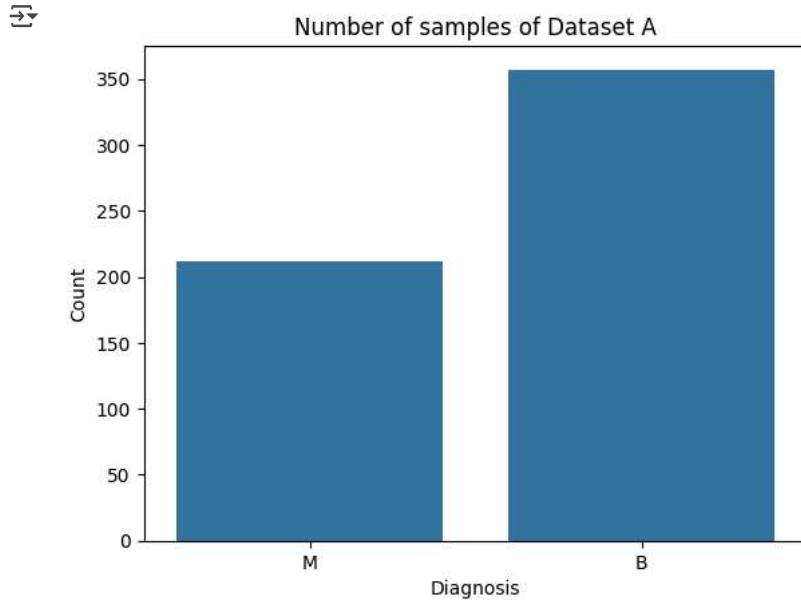
```
df.isnull().sum()
```

```
diagnosis          0
radius_mean        0
texture_mean       0
perimeter_mean    0
area_mean          0
smoothness_mean   0
compactness_mean  0
concavity_mean    0
concave points_mean 0
symmetry_mean     0
fractal_dimension_mean 0
radius_se          0
texture_se          0
perimeter_se       0
area_se            0
smoothness_se     0
compactness_se    0
concavity_se       0
concave points_se 0
symmetry_se        0
fractal_dimension_se 0
radius_worst       0
texture_worst      0
perimeter_worst   0
area_worst         0
smoothness_worst  0
compactness_worst 0
concavity_worst   0
concave points_worst 0
symmetry_worst    0
fractal_dimension_worst 0
```

```
print(df.diagnosis)
```

```
0      M
1      M
2      M
3      M
4      M
 ..
564    M
565    M
566    M
567    M
568    B
Name: diagnosis, Length: 569, dtype: object
```

```
import matplotlib.pyplot as plt
import seaborn as sns
sns.countplot(x='diagnosis', data=df)
plt.title('Number of samples of Dataset A')
plt.xlabel('Diagnosis')
plt.ylabel('Count')
plt.savefig('Number of samples of Dataset A.pdf', bbox_inches='tight')
plt.show()
```



```
df = df[df['diagnosis'] != 'ckd\t']
df.value_counts('diagnosis')
```

```
count
diagnosis
B    357
M    212
```

**dtype:** int64

```
x = df.drop(columns=['diagnosis'])
```

```
df.head()
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	symme
0	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	
1	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	
2	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	
3	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	
4	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	

5 rows × 32 columns

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
 #   Column            Non-Null Count  Dtype  
 ---  -- 
 0   diagnosis        569 non-null   object
```

```

1   radius_mean      569 non-null    float64
2   texture_mean     569 non-null    float64
3   perimeter_mean   569 non-null    float64
4   area_mean        569 non-null    float64
5   smoothness_mean  569 non-null    float64
6   compactness_mean 569 non-null    float64
7   concavity_mean   569 non-null    float64
8   concave points_mean 569 non-null    float64
9   symmetry_mean   569 non-null    float64
10  fractal_dimension_mean 569 non-null    float64
11  radius_se        569 non-null    float64
12  texture_se       569 non-null    float64
13  perimeter_se    569 non-null    float64
14  area_se          569 non-null    float64
15  smoothness_se   569 non-null    float64
16  compactness_se  569 non-null    float64
17  concavity_se    569 non-null    float64
18  concave points_se 569 non-null    float64
19  symmetry_se     569 non-null    float64
20  fractal_dimension_se 569 non-null    float64
21  radius_worst    569 non-null    float64
22  texture_worst   569 non-null    float64
23  perimeter_worst 569 non-null    float64
24  area_worst       569 non-null    float64
25  smoothness_worst 569 non-null    float64
26  compactness_worst 569 non-null    float64
27  concavity_worst 569 non-null    float64
28  concave points_worst 569 non-null    float64
29  symmetry_worst  569 non-null    float64
30  fractal_dimension_worst 569 non-null    float64
31  Unnamed: 32      0 non-null     float64
dtypes: float64(31), object(1)
memory usage: 142.4+ KB

```

```
df_encoded = pd.get_dummies(df, columns=['diagnosis'], drop_first=True)
```

```

# Convert 'bool' columns to int (True -> 1, False -> 0)
bool_cols = df.select_dtypes(include='bool').columns
df[bool_cols] = df[bool_cols].astype(int)

# Convert 'object' columns to numeric using Label Encoding
from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
object_cols = df.select_dtypes(include='object').columns

for col in object_cols:
    df[col] = df[col].astype(str) # Ensure all values are strings
    df[col] = label_encoder.fit_transform(df[col])

# Check the result
print(df.head())
print(df.dtypes)

diagnosis  radius_mean  texture_mean  perimeter_mean  area_mean  \
0           1         17.99       10.38      122.80     1001.0
1           1         20.57       17.77      132.90     1326.0
2           1         19.69       21.25      130.00     1203.0
3           1         11.42       20.38       77.58      386.1
4           1         20.29       14.34      135.10     1297.0

smoothness_mean  compactness_mean  concavity_mean  concave points_mean  \
0      0.11840        0.27760       0.3001       0.14710
1      0.08474        0.07864       0.0869       0.07017
2      0.10960        0.15990       0.1974       0.12790
3      0.14250        0.28390       0.2414       0.10520
4      0.10030        0.13280       0.1980       0.10430

symmetry_mean  ...  texture_worst  perimeter_worst  area_worst  \
0      0.2419  ...        17.33      184.60     2019.0
1      0.1812  ...        23.41      158.80     1956.0
2      0.2069  ...        25.53      152.50     1709.0
3      0.2597  ...        26.50       98.87      567.7
4      0.1809  ...        16.67      152.20     1575.0

smoothness_worst  compactness_worst  concavity_worst  concave points_worst  \
0      0.1622        0.6656       0.7119       0.2654
1      0.1238        0.1866       0.2416       0.1860
2      0.1444        0.4245       0.4504       0.2430
3      0.2098        0.8663       0.6869       0.2575
4      0.1374        0.2050       0.4000       0.1625

```

```

symmetry_worst fractal_dimension_worst Unnamed: 32
0      0.4601          0.11890      NaN
1      0.2750          0.08902      NaN
2      0.3613          0.08758      NaN
3      0.6638          0.17300      NaN
4      0.2364          0.07678      NaN

```

```

[5 rows x 32 columns]
diagnosis           int64
radius_mean         float64
texture_mean        float64
perimeter_mean     float64
area_mean           float64
smoothness_mean    float64
compactness_mean   float64
concavity_mean     float64
concave_points_mean float64
symmetry_mean      float64
fractal_dimension_mean float64
radius_se           float64
texture_se          float64
perimeter_se        float64
area_se             float64
smoothness_se       float64
compactness_se      float64
concavity_se        float64
concave_points_se   float64
symmetry_se         float64
fractal_dimension_se float64
radius_worst        float64

```

```
df.head()
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave_points_mean	symmetry_mean
0	1	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	
1	1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	
2	1	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	
3	1	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	
4	1	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	

```
5 rows × 32 columns
```

```
df.tail()
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave_points_mean	symmetry_mean
564	1	21.56	22.39	142.00	1479.0	0.11100	0.11590	0.24390	0.13890	
565	1	20.13	28.25	131.20	1261.0	0.09780	0.10340	0.14400	0.09791	
566	1	16.60	28.08	108.30	858.1	0.08455	0.10230	0.09251	0.05302	
567	1	20.60	29.33	140.10	1265.0	0.11780	0.27700	0.35140	0.15200	
568	0	7.76	24.54	47.92	181.0	0.05263	0.04362	0.00000	0.00000	

```
5 rows × 32 columns
```

```
num_cols = df.shape[1]
```

```
print("Number of columns:", num_cols)
```

```
Number of columns: 32
```

```
num_cols = len(df.columns) # ❌ This gives an integer
```

```
num_cols = df.select_dtypes(include=['int64', 'float64']).columns
```

```

for col in num_cols:
    print(col)
    print('Skew :', round(df[col].skew(), 2))

plt.figure(figsize=(15, 4))

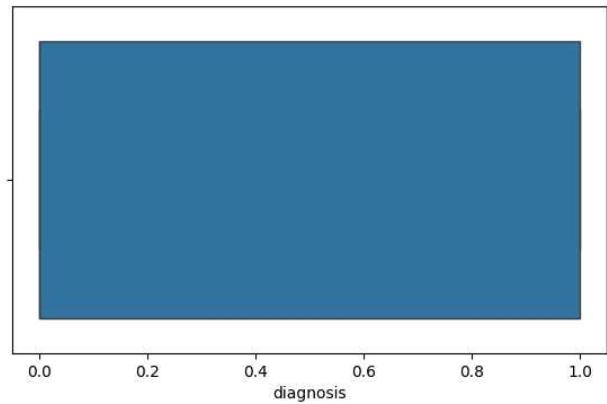
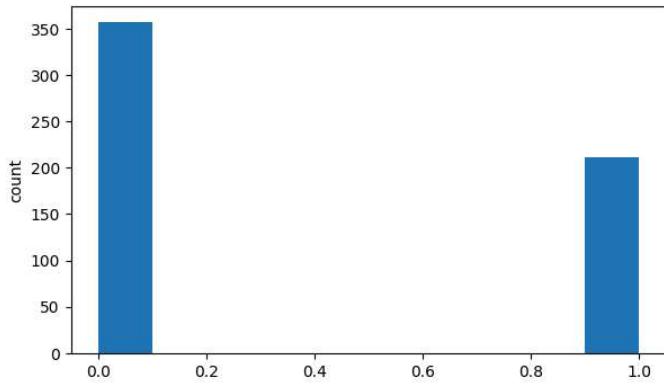
# Histogram
plt.subplot(1, 2, 1)
df[col].hist(grid=False)
plt.ylabel('count')

# Boxplot
plt.subplot(1, 2, 2)
sns.boxplot(x=df[col].dropna())

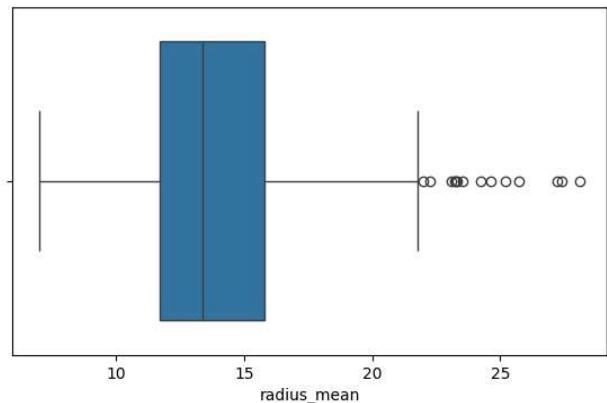
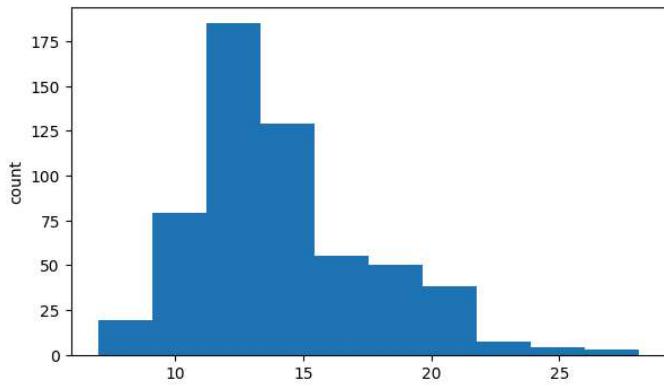
plt.show()

```

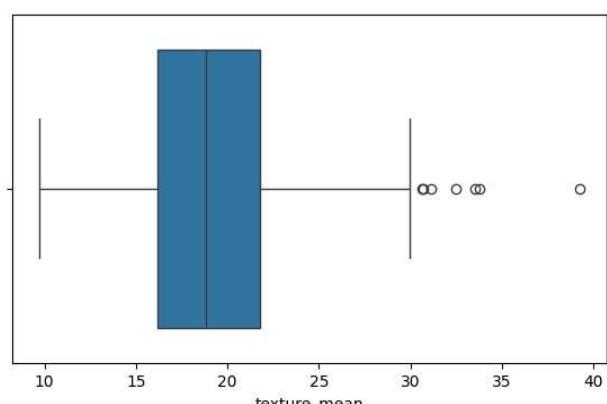
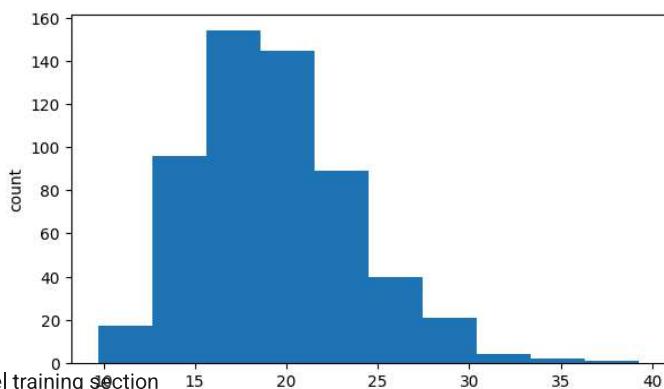
diagnosis  
Skew : 0.53



radius\_mean  
Skew : 0.94



texture\_mean  
Skew : 0.65



```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

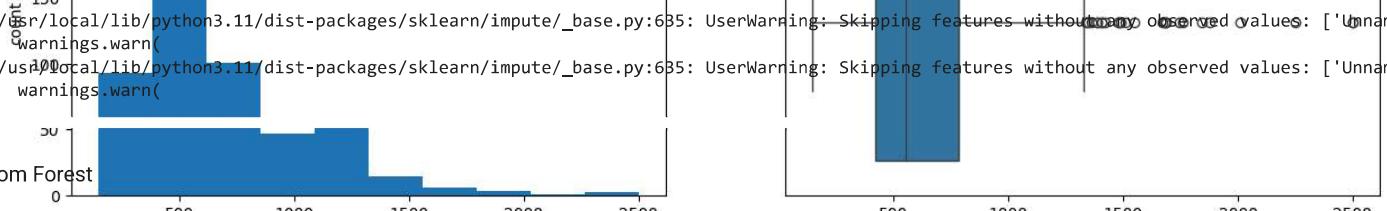
```
x = df.drop(columns=['diagnosis'])
y = df['diagnosis']

import numpy as np
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler

# Impute missing values using the mean
imputer = SimpleImputer(missing_values=np.nan, strategy='mean')
X_train = imputer.fit_transform(X_train)
X_test = imputer.transform(X_test)

# Feature normalization
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

→ /usr/local/lib/python3.11/dist-packages/sklearn/impute/\_base.py:635: UserWarning: Skipping features without any observed values: ['Unnamed: 32']  
 warnings.warn('/usr/local/lib/python3.11/dist-packages/sklearn/impute/\_base.py:635: UserWarning: Skipping features without any observed values: ['Unna  
 warnings.warn('Random Forest



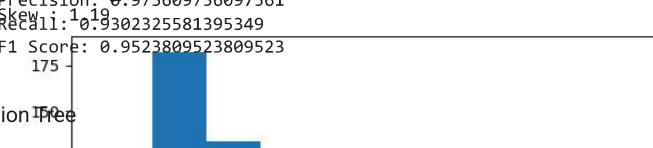
```
from sklearn.ensemble import RandomForestClassifier
rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train, y_train)
y_pred = rf_model.predict(X_test)
```

→ from sklearn.metrics import accuracy\_score, f1\_score, precision\_score, recall\_score  
# Calculate the metrics  
accuracy = accuracy\_score(y\_test, y\_pred)  
precision = precision\_score(y\_test, y\_pred)  
recall = recall\_score(y\_test, y\_pred)  
f1 = f1\_score(y\_test, y\_pred)

# Print the metrics  
print(f"Accuracy: {accuracy}")  
print(f"Precision: {precision}")  
print(f"Recall: {recall}")  
print(f"F1 Score: {f1}")

→ Accuracy: 0.9649122807017544  
Precision: 0.975609756097561  
Recall: 1.0  
F1 Score: 0.9523809523809523

Decision Tree



```
from sklearn.tree import DecisionTreeClassifier
dt_model = DecisionTreeClassifier(random_state=42)
dt_model.fit(X_train, y_train)
y_pred_dt = dt_model.predict(X_test)
```

→ from sklearn.metrics import accuracy\_score, f1\_score, precision\_score, recall\_score  
# Calculate the metrics  
accuracy\_dt = accuracy\_score(y\_test, y\_pred\_dt)  
precision\_dt = precision\_score(y\_test, y\_pred\_dt)  
recall\_dt = recall\_score(y\_test, y\_pred\_dt)  
f1\_dt = f1\_score(y\_test, y\_pred\_dt)

# Print the metrics  
print(f"Accuracy: {accuracy\_dt}")  
print(f"Precision: {precision\_dt}")  
print(f"Recall: {recall\_dt}")  
print(f"F1 Score: {f1\_dt}")

```
Accuracy: 0.9473684210526315
Precision: 0.9302325581395349
Recall: 0.9302325581395349
F1 Score: 0.9302325581395349
```

```
from sklearn.neighbors import KNeighborsClassifier
knn_model = KNeighborsClassifier()
knn_model.fit(X_train, y_train)
y_pred_knn = knn_model.predict(X_test)

# Print the metrics
print(f"Accuracy: {accuracy_knn}")
print(f"Precision: {precision_knn}")
print(f"Recall: {recall_knn}")
print(f"F1 Score: {f1_knn}")

# Calculate the metrics
accuracy_knn = accuracy_score(y_test, y_pred_knn)
precision_knn = precision_score(y_test, y_pred_knn)
recall_knn = recall_score(y_test, y_pred_knn)
f1_knn = f1_score(y_test, y_pred_knn)
```

```
Accuracy: 0.9473684210526315
Precision: 0.9302325581395349
Recall: 0.9302325581395349
F1 Score: 0.9302325581395349
symmetry_mean
Skew : 0.73
```

### Logistic Regression

```
from sklearn.linear_model import LogisticRegression
lr_model = LogisticRegression()
lr_model.fit(X_train, y_train)
y_pred_lr = lr_model.predict(X_test)

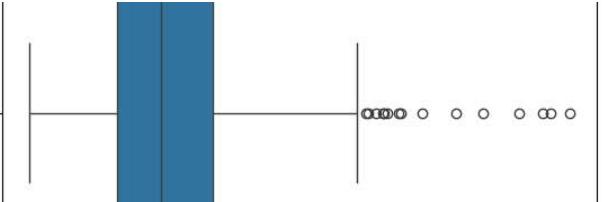
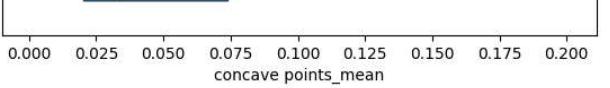
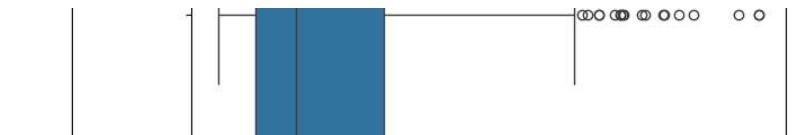
# Print the metrics
print(f"Accuracy: {accuracy_lr}")
print(f"Precision: {precision_lr}")
print(f"Recall: {recall_lr}")
print(f"F1 Score: {f1_lr}")

# Calculate the metrics
accuracy_lr = accuracy_score(y_test, y_pred_lr)
precision_lr = precision_score(y_test, y_pred_lr)
recall_lr = recall_score(y_test, y_pred_lr)
f1_lr = f1_score(y_test, y_pred_lr)
```

```
Accuracy: 0.9736842105263158
Precision: 0.9761904761904762
Recall: 0.9534883720930233
F1 Score: 0.9647058823529412
```

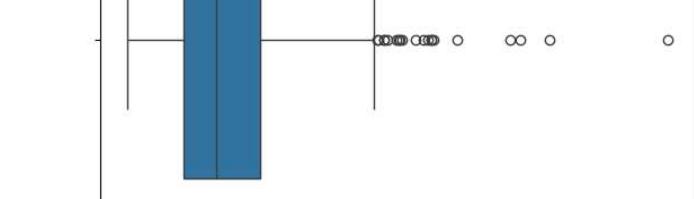
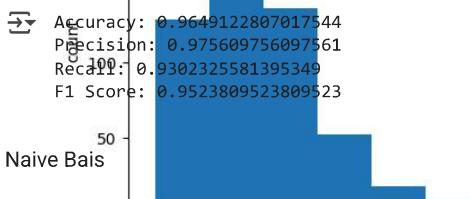
### Adaboost

```
from sklearn.ensemble import AdaBoostClassifier
adaboost_model = AdaBoostClassifier(random_state=42)
adaboost_model.fit(X_train, y_train)
y_pred_ada = adaboost_model.predict(X_test)
```



```
from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
# Calculate the metrics
accuracy_ada = accuracy_score(y_test, y_pred_ada)
precision_ada = precision_score(y_test, y_pred_ada)
recall_ada = recall_score(y_test, y_pred_ada)
f1_ada = f1_score(y_test, y_pred_ada)

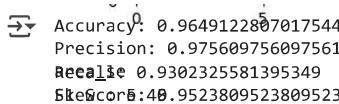
# Print the metrics
print(f"Accuracy: {accuracy_ada}")
print(f"Precision: {precision_ada}")
print(f"Recall: {recall_ada}")
print(f"F1 Score: {f1_ada}")
```



```
from sklearn.naive_bayes import GaussianNB
nb_model = GaussianNB()
nb_model.fit(X_train, y_train)
y_pred_nb = nb_model.predict(X_test)
```

```
# Calculate the metrics
accuracy_nb = accuracy_score(y_test, y_pred_nb)
precision_nb = precision_score(y_test, y_pred_nb)
recall_nb = recall_score(y_test, y_pred_nb)
f1_nb = f1_score(y_test, y_pred_nb)
```

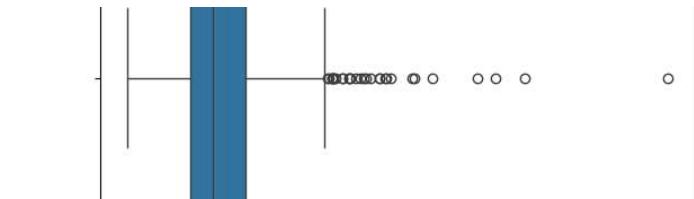
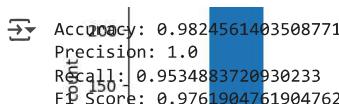
```
# Print the metrics
print(f"Accuracy: {accuracy_nb}")
print(f"Precision: {precision_nb}")
print(f"Recall: {recall_nb}")
print(f"F1 Score: {f1_nb}")
```



```
from sklearn.svm import SVC
svm_model = SVC(random_state=42)
svm_model.fit(X_train, y_train)
y_pred_svm = svm_model.predict(X_test)
```

```
# Calculate the metrics
accuracy_svm = accuracy_score(y_test, y_pred_svm)
precision_svm = precision_score(y_test, y_pred_svm)
recall_svm = recall_score(y_test, y_pred_svm)
f1_svm = f1_score(y_test, y_pred_svm)
```

```
# Print the metrics
print(f"Accuracy: {accuracy_svm}")
print(f"Precision: {precision_svm}")
print(f"Recall: {recall_svm}")
print(f"F1 Score: {f1_svm}")
```



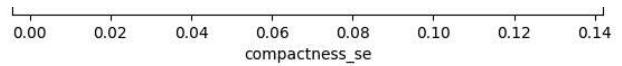
```
from xgboost import XGBClassifier
xgb_model = XGBClassifier(random_state=42)
```

```
xgboost_model.fit(X_train, y_train)
y_pred_xgb = xgboost_model.predict(X_test)

from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
# Calculate the metrics
accuracy_xgb= accuracy_score(y_test, y_pred_xgb)
precision_xgb = precision_score(y_test, y_pred_xgb)
recall_xgb= recall_score(y_test, y_pred_xgb)
f1_xgb= f1_score(y_test, y_pred_xgb)

# Print the metrics
print(f"Accuracy: {accuracy_xgb}")
print(f"Precision: {precision_xgb}")
print(f"Recall: {recall_xgb}")
print(f"F1 Score: {f1_xgb}")
```

→ Accuracy: 0.956140350877193  
 Precision: 0.9523809523809523  
 Recall: 0.9302325581395349  
 F1 score: 0.9411764705882353  
 Skew: 5.11



```
from keras.models import Sequential
from keras.layers import Dense, Dropout # Add Dropout import
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
# Define ANN model
ann_model = Sequential()
ann_model.add(Dense(64, input_dim=X_train.shape[1], activation='relu'))
ann_model.add(Dropout(0.2))
ann_model.add(Dense(32, activation='relu'))
ann_model.add(Dense(1, activation='sigmoid'))
```

```
# Compile the model
ann_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
# Train the model
ann_model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_test, y_test))
```

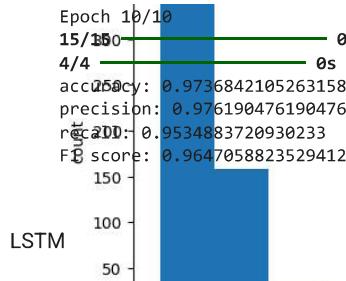
```
# Evaluate the model
ann_predictions = ann_model.predict(X_test)
y_pred_ann = (ann_predictions > 0.5).astype("int32") # Correct variable name
```

```
# Calculate evaluation metrics
accuracy = accuracy_score(y_test, y_pred_ann)
precision = precision_score(y_test, y_pred_ann)
recall = recall_score(y_test, y_pred_ann)
f1 = f1_score(y_test, y_pred_ann)
```

```
# Output results
print(f"accuracy: {accuracy}")
print(f"precision: {precision}")
print(f"recall: {recall}")
print(f"F1 score: {f1}")
```

→ /usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:93: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to `super().\_\_init\_\_`. Activity regularization is handled by the `activity\_regularizer` argument.

Epoch 1/10  
 15/15 4s 70ms/step - accuracy: 0.5680 - loss: 0.6547 - val\_accuracy: 0.9386 - val\_loss: 0.3926  
 Epoch 2/10  
 15/15 0s 20ms/step - accuracy: 0.9242 - loss: 0.3794 - val\_accuracy: 0.9561 - val\_loss: 0.2380  
 Epoch 3/10  
 15/15 1s 27ms/step - accuracy: 0.9367 - loss: 0.2384 - val\_accuracy: 0.9649 - val\_loss: 0.1533  
 Epoch 4/10  
 15/15 1s 66ms/step - accuracy: 0.9401 - loss: 0.1835 - val\_accuracy: 0.9649 - val\_loss: 0.1138  
 Epoch 5/10  
 15/15 1s 46ms/step - accuracy: 0.9587 - loss: 0.1319 - val\_accuracy: 0.9649 - val\_loss: 0.0934  
 Epoch 6/10  
 15/15 0s 20ms/step - accuracy: 0.9721 - loss: 0.1114 - val\_accuracy: 0.9825 - val\_loss: 0.0820  
 Epoch 7/10 0.01 0.02 0.03 0.04 0.05 0.06 0.07 0.08  
 15/15 0s 12ms/step - accuracy: 0.9484 - loss: 0.1106 - val\_accuracy: 0.9737 - val\_loss: 0.0753  
 Epoch 8/10  
 15/15 0s 13ms/step - accuracy: 0.9809 - loss: 0.0809 - val\_accuracy: 0.9825 - val\_loss: 0.0704  
 Epoch 9/10  
 15/15 0s 11ms/step - accuracy: 0.9718 - loss: 0.0863 - val\_accuracy: 0.9825 - val\_loss: 0.0674  
 15/15 350 1



```
import numpy as np
import tensorflow as tf
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dropout, Dense
from sklearn.model_selection import train_test_split
```

```
if len(X_train.shape) == 2:
    X_train = np.array(X_train).reshape((X_train.shape[0], 1, X_train.shape[1]))
    X_test = np.array(X_test).reshape((X_test.shape[0], 1, X_test.shape[1]))
```

```
# Build LSTM model
model_lstm = Sequential()
model_lstm.add(LSTM(150, activation='relu', return_sequences=True, input_shape=(X_train.shape[1], X_train.shape[2])))
model_lstm.add(Dropout(0.2))
model_lstm.add(LSTM(150, activation='relu'))
model_lstm.add(Dense(1, activation='sigmoid'))
```

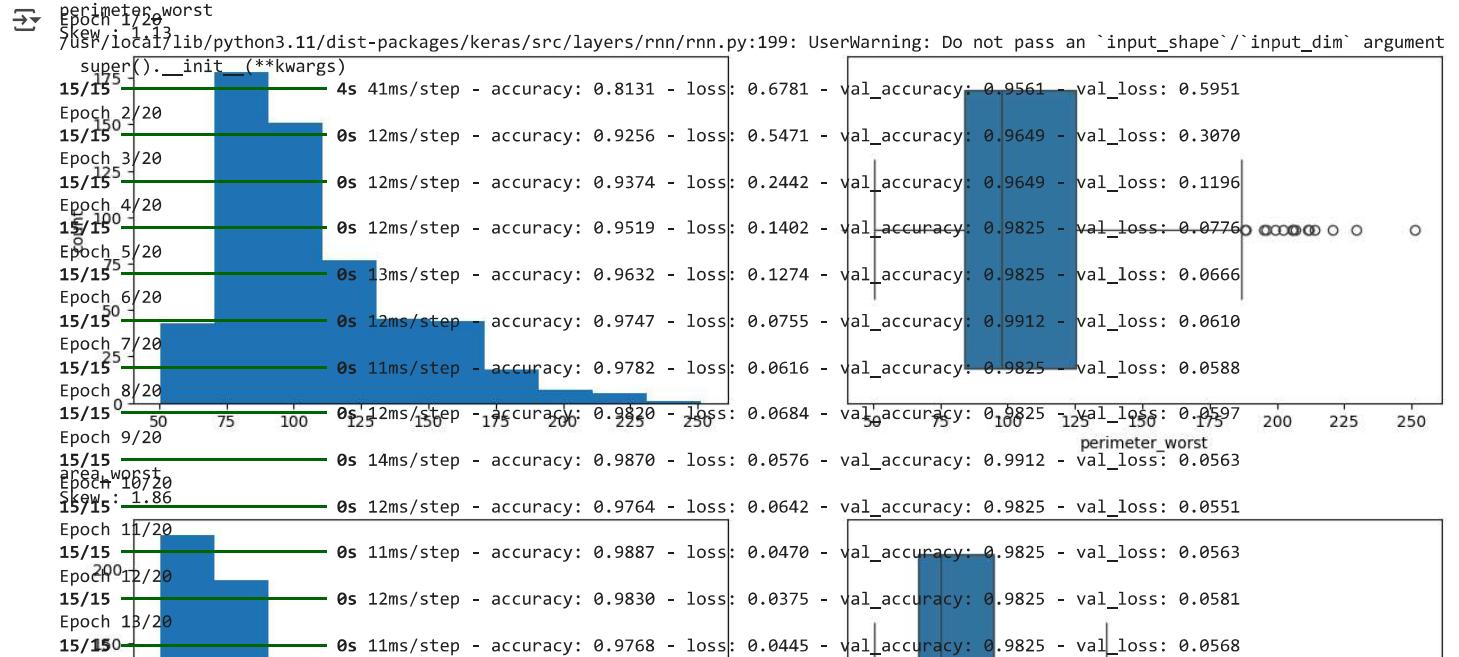
```
# Compile the model
model_lstm.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

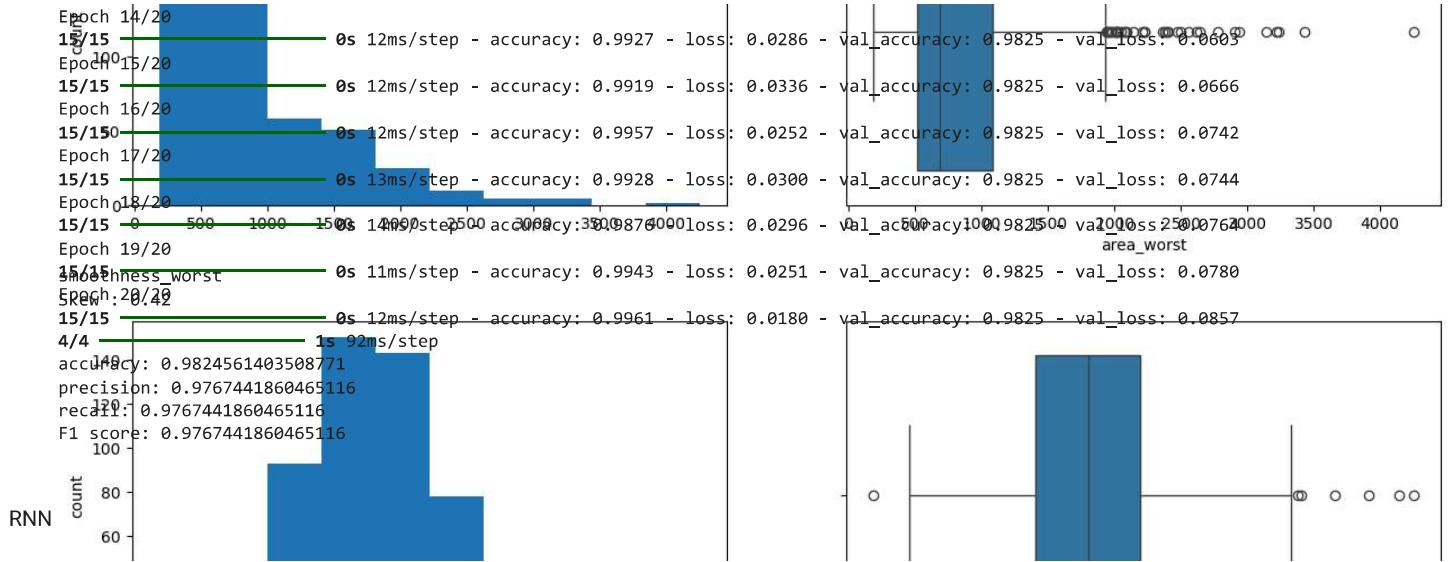
```
# Train the model
model_lstm.fit(X_train, y_train, epochs=20, batch_size=32, validation_data=(X_test, y_test))
```

```
# Make predictions
y_pred = model_lstm.predict(X_test)
y_pred_lstm = (y_pred > 0.5).astype("int32")
```

```
# Calculate accuracy, precision, recall, and F1-score
accuracy = accuracy_score(y_test, y_pred_lstm)
precision = precision_score(y_test, y_pred_lstm)
recall = recall_score(y_test, y_pred_lstm)
f1 = f1_score(y_test, y_pred_lstm)
```

```
# Output results
print(f"accuracy: {accuracy}")
print(f"precision: {precision}")
print(f"recall: {recall}")
print(f"F1 score: {f1}")
```





```

import matplotlib.pyplot as plt
import numpy as np

# Model names and their corresponding metrics
models = ['RF', 'DT', 'KNN', 'LR', 'ADA', 'NB', 'SVM', 'XGB', 'ANN', 'LSTM', 'RNN']

# Accuracy, Precision, Recall, F1 Score for each model (updated based on the provided data)
accuracy = [100 * x for x in [
    0.9307228915662651, 0.8855421686746988, 0.9186746987951807, 0.9367469879518072, 0.9246987951807228,
    0.927710843373494, 0.927710843373494, 0.9397590361445783, 0.9307228915662651, 0.9246987951807228,
    0.9216867469879518
]]
precision = [100 * x for x in [
    0.9305135951661632, 0.9470198675496688, 0.9270516717325228, 0.9415384615384615, 0.9435736677115988,
    0.9382716049382716, 0.927710843373494, 0.9444444444444444, 0.9305135951661632, 0.9353846153846154,
    0.92727272727272
]]
recall = [100 * x for x in [
    1.0, 0.9285714285714286, 0.9902597402597403, 0.9935064935064936, 0.9772727272727273,
    0.987012987012987, 1.0, 0.9935064935064936, 1.0, 0.987012987012987, 0.9935064935064936
]]
f1_score = [100 * x for x in [
    0.9640062597809077, 0.9377049180327869, 0.957613814756672, 0.966824644549763, 0.960127591706539,
    0.9620253164556962, 0.9625, 0.9683544303797469, 0.9640062597809077, 0.9605055292259084,
    0.9592476489028213
]]

# Dummy error values for error bars (calculated or assumed for the sake of this example)
error_accuracies = [0.02, 0.03, 0.02, 0.02, 0.02, 0.03, 0.02, 0.01, 0.02, 0.02, 0.05]
error_precisions = [0.01, 0.02, 0.01, 0.02, 0.02, 0.02, 0.01, 0.01, 0.01, 0.02, 0.05]
error_recalls = [0.01, 0.01, 0.02, 0.01, 0.01, 0.02, 0.01, 0.01, 0.01, 0.01, 0.05]
error_f1_scores = [0.02, 0.03, 0.02, 0.02, 0.02, 0.03, 0.02, 0.01, 0.02, 0.02, 0.05]

# Metrics and labels
metrics = [accuracy, f1_score, precision, recall]
metrics_labels = ['Accuracy', 'F1 Score', 'Precision', 'Recall']
metrics_errors = [error_accuracies, error_f1_scores, error_precisions, error_recalls]

# Bar chart setup
x_pos = np.arange(len(models))
width = 0.2 # Bar width

fig, ax = plt.subplots(figsize=(30, 15))

# Plot grouped bars with error bars
for i, (metric, label, errors) in enumerate(zip(metrics, metrics_labels, metrics_errors)):
    bar = ax.bar(
        x_pos + i * width, metric, width * 0.8, yerr=errors,
        label=label, alpha=0.7, ecolor='black', capsized=5
    )

    # Add labels for the bar values inside the bars
    for j, (rect, value) in enumerate(zip(bar, metric)):
        height = rect.get_height()

```

```

offset_x = rect.get_x() + rect.get_width() / 2.0
offset_y = height * 0.5

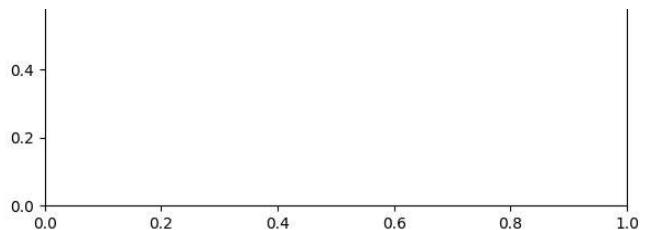
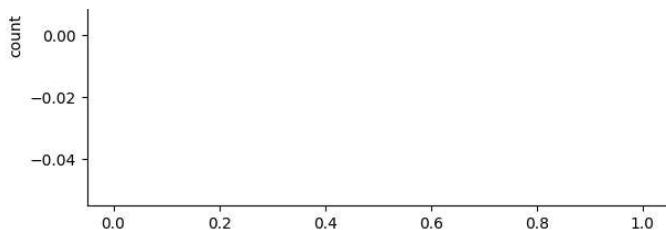
# For `0.0%` values, place the text inside the bar with a small gap
if value == 0.0:
    ax.text(
        offset_x,
        1, # Add a small gap above the bottom of the bar
        f'{value:.1f}%',
        ha='center',
        va='center',
        fontsize=18,
        color='black',
        rotation=5, # No rotation for 0.0%
        rotation_mode='anchor'
    )

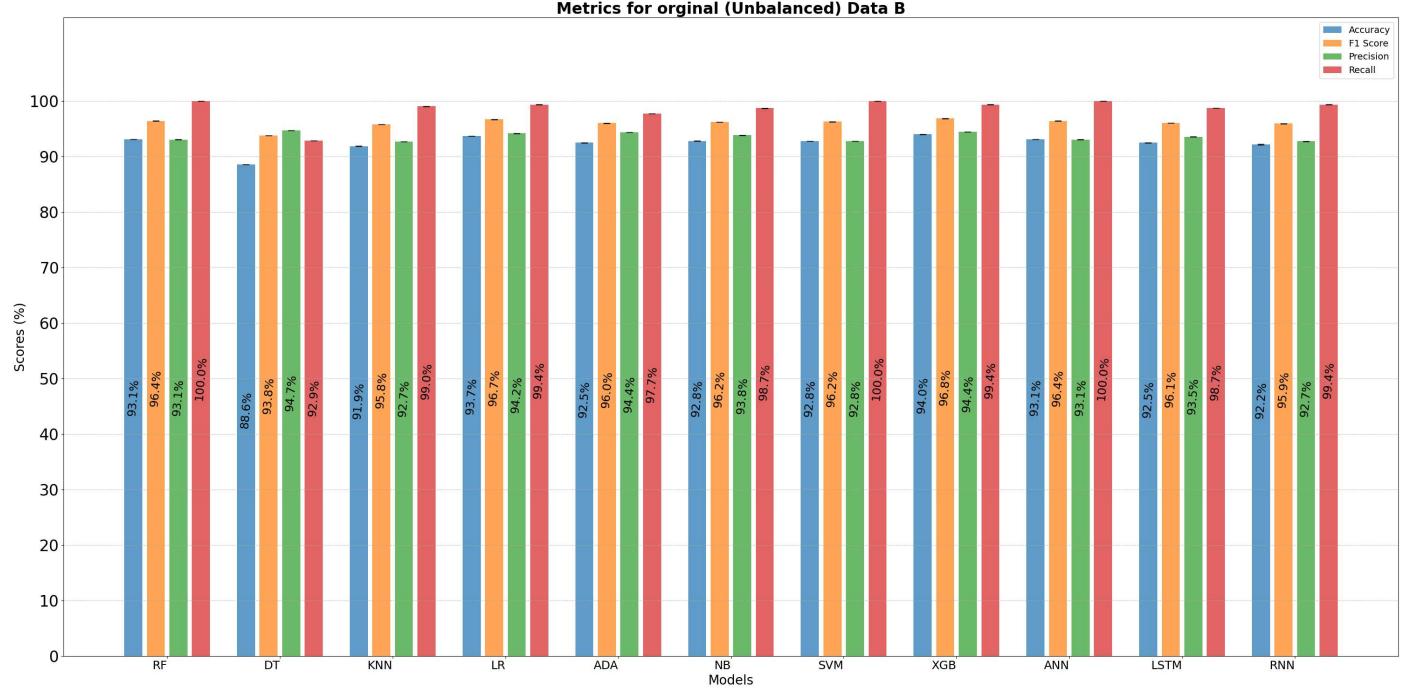
else:
    ax.text(
        offset_x,
        offset_y,
        f'{value:.1f}%',
        ha='center',
        va='center',
        fontsize=18,
        color='black',
        rotation=85, # Rotate the non-zero text
        rotation_mode='anchor'
    )
x_pos_adjusted = x_pos.copy()
x_pos_adjusted[-1] += 0.2

# Styling and layout
ax.set_title('Metrics for orginal (Unbalanced) Data B', fontsize=24, weight='bold')
ax.set_xlabel('Models', fontsize=20)
ax.set_ylabel('Scores (%)', fontsize=20)
ax.set_xticks(x_pos + width * 1.5)
ax.set_xticklabels(models, ha='right', fontsize=18)
ax.set_xlim(0, 115)
ax.set_yticks(np.arange(0, 110, 10))
ax.tick_params(axis='y', labelsize=24)
ax.legend(fontsize=14, loc='upper right', bbox_to_anchor=(1, 1))
ax.yaxis.grid(True, linestyle='--', alpha=0.7)

# Adjust layout and save the figure
plt.tight_layout()
plt.savefig('Metrics for orginal (Unbalanced) Data B.pdf', bbox_inches='tight')
plt.show()

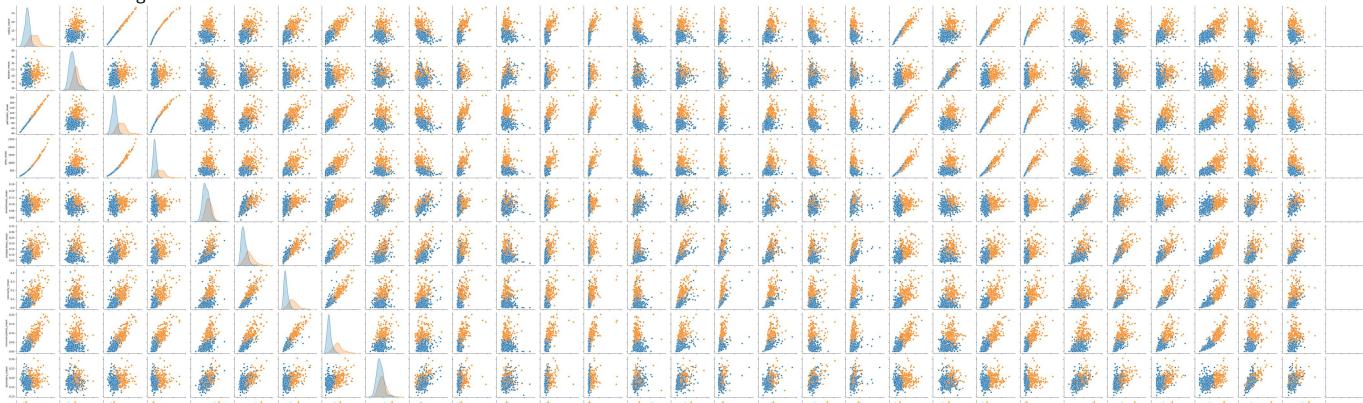
```





```
df.info()
sns.pairplot(df, hue = 'diagnosis')
```

```
↳ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   diagnosis        569 non-null    int64  
 1   radius_mean      569 non-null    float64 
 2   texture_mean     569 non-null    float64 
 3   perimeter_mean   569 non-null    float64 
 4   area_mean        569 non-null    float64 
 5   smoothness_mean  569 non-null    float64 
 6   compactness_mean 569 non-null    float64 
 7   concavity_mean   569 non-null    float64 
 8   concave points_mean 569 non-null    float64 
 9   symmetry_mean    569 non-null    float64 
 10  fractal_dimension_mean 569 non-null    float64 
 11  radius_se        569 non-null    float64 
 12  texture_se       569 non-null    float64 
 13  perimeter_se    569 non-null    float64 
 14  area_se          569 non-null    float64 
 15  smoothness_se   569 non-null    float64 
 16  compactness_se  569 non-null    float64 
 17  concavity_se    569 non-null    float64 
 18  concave points_se 569 non-null    float64 
 19  symmetry_se     569 non-null    float64 
 20  fractal_dimension_se 569 non-null    float64 
 21  radius_worst    569 non-null    float64 
 22  texture_worst   569 non-null    float64 
 23  perimeter_worst 569 non-null    float64 
 24  area_worst       569 non-null    float64 
 25  smoothness_worst 569 non-null    float64 
 26  compactness_worst 569 non-null    float64 
 27  concavity_worst 569 non-null    float64 
 28  concave points_worst 569 non-null    float64 
 29  symmetry_worst  569 non-null    float64 
 30  fractal_dimension_worst 569 non-null    float64 
 31  Unnamed: 32      0 non-null     float64 
dtypes: float64(31), int64(1)
memory usage: 142.4 KB
<seaborn.axisgrid.PairGrid at 0x7dd8d736d10>
```



```
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np

# Compute the correlation matrix
corr = df.corr(numeric_only=False)

# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr, dtype=bool))

# Set up the matplotlib figure
plt.figure(figsize=(21, 21), dpi=700)

# Draw the heatmap
ax = sns.heatmap(corr, vmin=-0.1, vmax=1.0, mask=mask,
                  annot=True, fmt=".2f", cmap='coolwarm',
                  square=True, linewidths=0.9,
                  cbar_kws={"shrink":0.6, "pad":0.01})

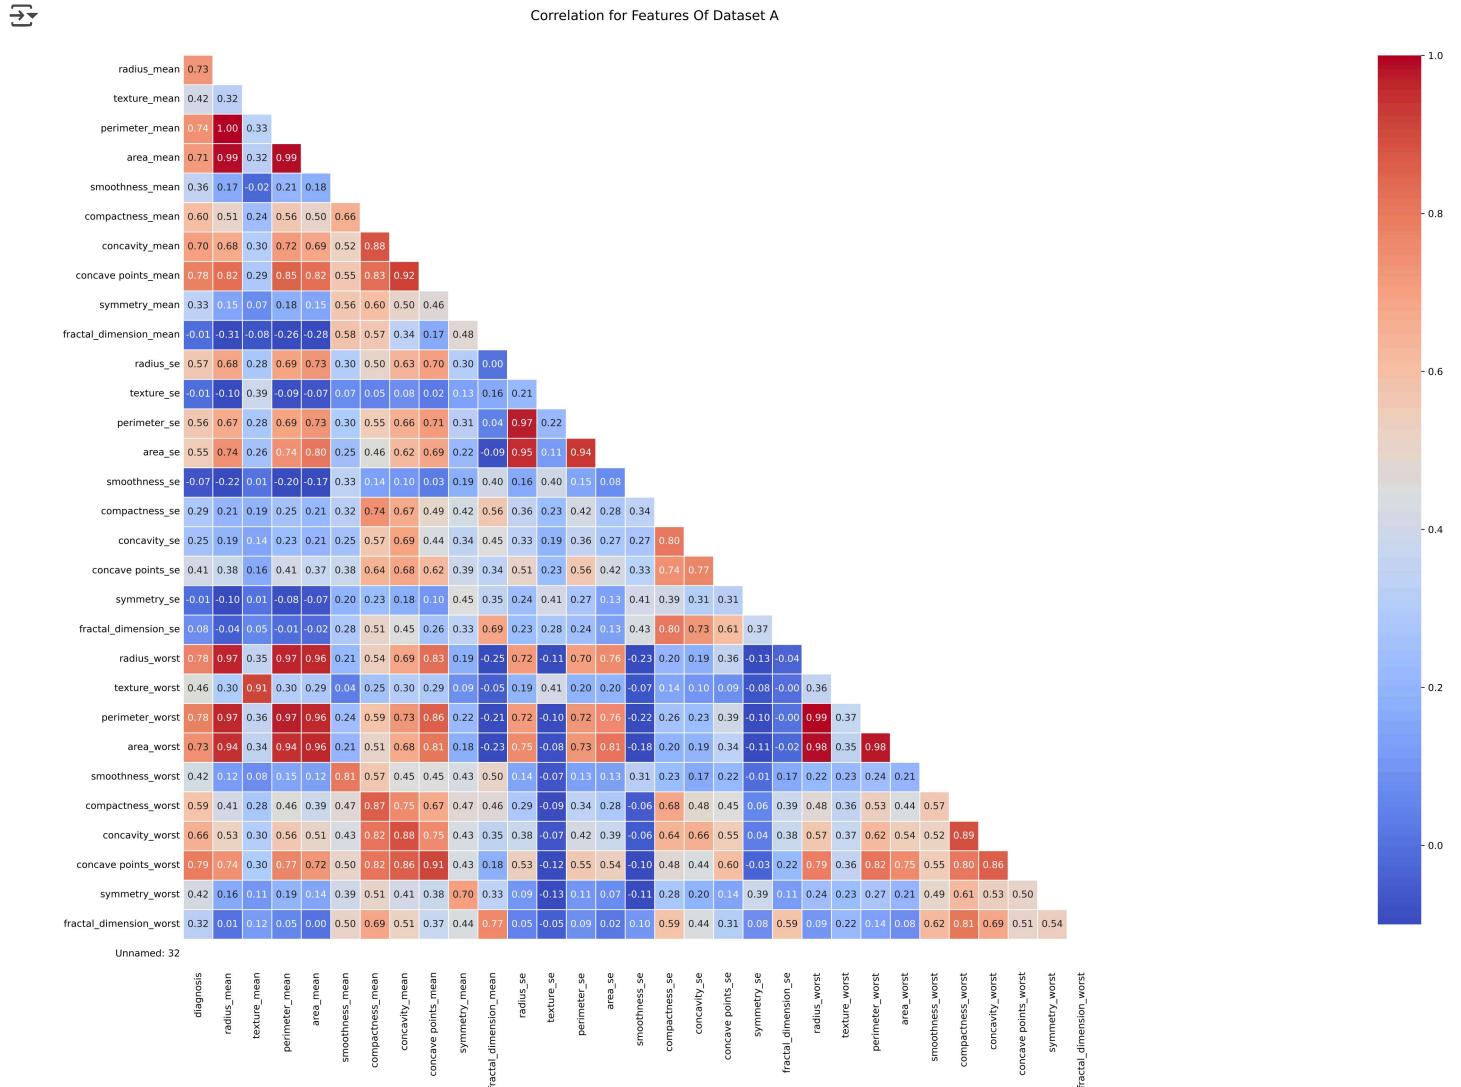
# Access the colorbar from the Axes object
cbar = ax.collections[0].colorbar
cbar.ax.set_position([0.95, 0.1, 0.03, 0.8])
```

```
# Customize labels
cust_labels = df.columns.tolist()
cust_labels[0] = ''
ax.set_yticklabels(cust_labels)
ax.tick_params(axis='y', which='both', length=0)
cust_labels = df.columns.tolist()
cust_labels[-1] = ''
ax.set_xticklabels(cust_labels)
ax.tick_params(axis='x', which='both', length=0)

plt.title("Correlation for Features Of Dataset A", fontsize=14)

# Save figure
plt.savefig('Correlation for Features Of Dataset A.pdf', bbox_inches='tight')

plt.show()
```



## Gain Ratio

```
X_res = df.drop(columns=['diagnosis'])
y_res = df['diagnosis']
```

```

import pandas as pd
import numpy as np
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import log_loss

# Compute Entropy of a distribution
def entropy(y_res):
    value_counts = y_res.value_counts(normalize=True)
    return -np.sum(value_counts * np.log2(value_counts + 1e-9))

# Compute Information Gain
def information_gain(X_res, y_res, feature):
    base_entropy = entropy(y_res)
    feature_entropy = 0
    for value in X_res[feature].unique():
        subset = y_res[X_res[feature] == value]
        feature_entropy += len(subset) / len(X_res) * entropy(subset)
    return base_entropy - feature_entropy

# Compute Split Information==
def split_information(X_res, feature):
    value_counts = X_res[feature].value_counts(normalize=True)
    return -np.sum(value_counts * np.log2(value_counts + 1e-9))

# Compute Gain Ratio
def gain_ratio(X_res, y_res, feature):
    ig = information_gain(X_res, y_res, feature)
    si = split_information(X_res, feature)
    return ig / (si + 1e-9)

# Compute Gain Ratios for all features
gain_ratios = {feature: gain_ratio(X_res, y_res, feature) for feature in X_res.columns}

# Convert to a DataFrame for easy visualization
gain_ratio_results = pd.DataFrame(gain_ratios.items(), columns=['Feature', 'Gain Ratio'])
gain_ratio_results = gain_ratio_results.sort_values(by="Gain Ratio", ascending=False)

# Display the results
print(gain_ratio_results)

```

	Feature	Gain Ratio
30	Unnamed: 32	9.526351e+08
7	concave points_mean	1.045177e-01
6	concavity_mean	1.039405e-01
27	concave points_worst	1.036364e-01
26	concavity_worst	1.033368e-01
16	concavity_se	1.033053e-01
20	radius_worst	1.032616e-01
12	perimeter_se	1.032500e-01
10	radius_se	1.032013e-01
2	perimeter_mean	1.031608e-01
23	area_worst	1.031573e-01
14	smoothness_se	1.030375e-01
13	area_se	1.027870e-01
3	area_mean	1.025957e-01
15	compactness_se	1.020111e-01
25	compactness_worst	1.019363e-01
19	fractal_dimension_se	1.011792e-01
5	compactness_mean	1.005991e-01
22	perimeter_worst	1.003446e-01
29	fractal_dimension_worst	9.871665e-02
0	radius_mean	9.856645e-02
17	concave points_se	9.660127e-02
21	texture_worst	9.651886e-02
11	texture_se	9.634091e-02
28	symmetry_worst	9.496908e-02
1	texture_mean	9.470897e-02
9	fractal_dimension_mean	9.395664e-02
18	symmetry_se	9.204348e-02
4	smoothness_mean	8.840055e-02
8	symmetry_mean	8.523176e-02
24	smoothness_worst	8.426232e-02

```

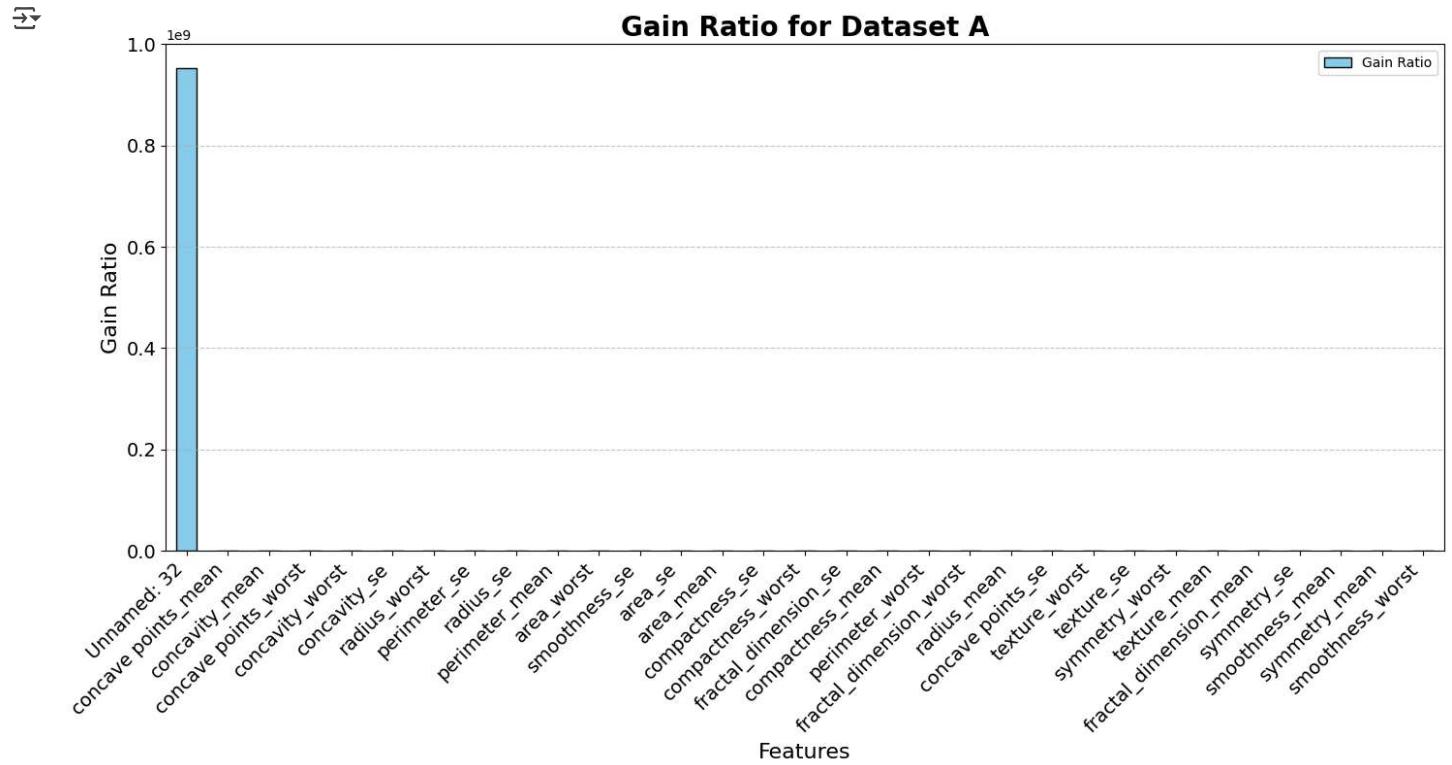
# Plotting the gain ratios
plt.figure(figsize=(28, 8))
ax1 = plt.subplot(1, 2, 1)
gain_ratio_results.plot(kind='bar', color='skyblue', edgecolor='black', ax=ax1)
ax1.set_title('Gain Ratio for Dataset A', fontsize=20, fontweight='bold')
ax1.set_xlabel('Features', fontsize=16)

```

```

ax1.set_ylabel('Gain Ratio', fontsize=16)
ax1.set_xticklabels(gain_ratio_results['Feature'], rotation=45, ha='right', fontsize=14)
ax1.tick_params(axis='y', labelsize=14) # Correct way to set y-tick label size
ax1.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.savefig("Gain Ratio for Dataset A.pdf") # save figure
plt.show()

```



## Information Gain

```

import pandas as pd
import numpy as np
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import log_loss

# Compute Entropy of a distribution
def entropy(y_res):
    value_counts = y_res.value_counts(normalize=True)
    return -np.sum(value_counts * np.log2(value_counts + 1e-9))

# Compute Information Gain
def information_gain(X_res, y_res, feature):
    base_entropy = entropy(y_res)
    feature_entropy = 0
    for value in X_res[feature].unique():
        subset = y_res[X_res[feature] == value]
        feature_entropy += len(subset) / len(X_res) * entropy(subset)
    return base_entropy - feature_entropy

# Compute Information Gains for all features
information_gains = {feature: information_gain(X_res, y_res, feature) for feature in X_res.columns}

# Convert to a DataFrame for easy visualization
information_gain_results = pd.DataFrame(information_gains.items(), columns=['Feature', 'Information Gain'])
information_gain_results = information_gain_results.sort_values(by="Information Gain", ascending=False)

# Display the results

```

```
print(information_gain_results)
```

	Feature	Information Gain
30	Unnamed: 32	0.952635
7	concave points_mean	0.942090
11	smoothness_se	0.925060